

Magnitude of the adaptive parameter step · all 24 matrices

AdamW adaptive update Muon post-NS update



$||LR \times \text{adaptive direction}||_F$, before weight decay and parameter-write rounding. Group learning rates differ.